


Recommendation stability and quality on the full evaluation suite

Gittins observations shared within each run · 1%-10% budget · Lower is better



Means of per-run metrics; bars: pointwise 95% paired seed-block bootstrap intervals (10,000 draws).
MMLU averages fixed subjects within each seed. Compare rules within a dataset; raw movement depends on update frequency.
MMLU: accuracy percentage points. AlpacaEval: score points on a 0-100 scale.