


Does LCB improve recommendations on other datasets?

Same observations for both rules. Lower regret is better. Budgets shown: 1–10%.



LCB strength fixed at $z = 1.645$. Bands: pointwise 95% paired seed-block bootstrap intervals (1,000 draws).

PIQA: 5 fixed matrices $\times$ 20 seeds per prior; each MMLU subject: 1 fixed matrix $\times$ 20 seeds per prior. No dataset resampling.