


Recommendation stability across the saved evaluation paths

1%–10% budget · Every actual recommendation update · Lower is better



Points: mean of per-run metrics; bars: pointwise 95% seed-block bootstrap intervals (10,000 draws).
Compare methods within each dataset: matrix sizes and update frequencies differ. Exploratory analysis on fixed matrices.